

Controlling Intelligent Agents in Virtual World

The OpenCog Embodiment Component

Nil Geisweiller

Novamente LLC

Xiamen University
AGI Summer School 2009

- 1 Introduction
- 2 Embodiment Sub-components
 - Proxy
 - Operational Agent (Pet) Controller
 - Learning Server
 - Collective Experience Store
- 3 Demo...
- 4 Conclusion

Outline

- 1 Introduction
- 2 Embodiment Sub-components
 - Proxy
 - Operational Agent (Pet) Controller
 - Learning Server
 - Collective Experience Store
- 3 Demo...
- 4 Conclusion

What is OpenCog Embodiment Component?

Embodiment component (formerly called PetBrain)

Interface OpenCog for virtual world agents or robots



Embodiment Component Function

Goals:

- Environment where to teach and test OpenCog

Embodiment Component Function

Goals:


- Environment where to teach and test OpenCog
- Virtual World \Rightarrow Easier than real world

Embodiment Component Function


Goals:

- Environment where to teach and test OpenCog
- Virtual World \Rightarrow Easier than real world
- But robot interface as well (in development)

Embodiment Component Overview




Embodiment Component Overview




- 1 Proxy, interface between the virtual world and Embodiment

Embodiment Component Overview




- 1 Proxy, interface between the virtual world and Embodiment
- 2 Operational Agent (Pet) Controller (OPC), reactive brain

Embodiment Component Overview




- 1 Proxy, **interface between the virtual world and Embodiment**
- 2 Operational Agent (Pet) Controller (OPC), **reactive brain**
- 3 Learning Server (LS), **imitation learning for now**

Embodiment Component Overview



- 1 Proxy, **interface between the virtual world and Embodiment**
- 2 Operational Agent (Pet) Controller (OPC), **reactive brain**
- 3 Learning Server (LS), **imitation learning for now**
- 4 Collective Experience Store (CES) (not implemented yet)

Embodiment Component Overview



- 1 Proxy, **interface between the virtual world and Embodiment**
- 2 Operational Agent (Pet) Controller (OPC), **reactive brain**
- 3 Learning Server (LS), **imitation learning for now**
- 4 Collective Experience Store (CES) (not implemented yet)
- 5 Router: communication with Sockets (distributed or not)


Outline

- 1 Introduction
- 2 Embodiment Sub-components**
 - Proxy
 - Operational Agent (Pet) Controller
 - Learning Server
 - Collective Experience Store
- 3 Demo...
- 4 Conclusion

Outline

- 1 Introduction
- 2 **Embodiment Sub-components**
 - **Proxy**
 - Operational Agent (Pet) Controller
 - Learning Server
 - Collective Experience Store
- 3 Demo...
- 4 Conclusion

Embodiment Proxy



Proxy main function

Abstract away all details of interacting with a Virtual World

For the moment Embodiment Proxy for

- Multiverse (implemented)



For the moment Embodiment Proxy for

- Multiverse (implemented)



- realXtend (almost complete)



For the moment Embodiment Proxy for

- Multiverse (implemented)
- realXtend (almost complete)
- Nao Robot (in development)



Embodiment Proxy sub-functions

Proxy handles:

- Input stream of **perceptual data**

Embodiment Proxy sub-functions

Proxy handles:

- Input stream of **perceptual data**
- Output stream of **actions and action feedback** (failure or success)

Embodiment Proxy sub-functions

Proxy handles:


- Input stream of **perceptual data**
- Output stream of **actions and action feedback** (failure or success)
- Maintaining a Global World Map (GWM)
 - Shared GWM for all agents
 - ⇒ reduce message traffic agents ↔ virtual world

Embodiment Proxy sub-functions

Proxy handles:

- Input stream of **perceptual data**
- Output stream of **actions and action feedback** (failure or success)
- Maintaining a Global World Map (GWM)
 - Shared GWM for all agents
⇒ reduce message traffic agents ↔ virtual world
- Commands from a human avatar
 - messages to the agents, order, queries
 - meta-commands (load agent, etc)

Embodiment Proxy ↔ PAI

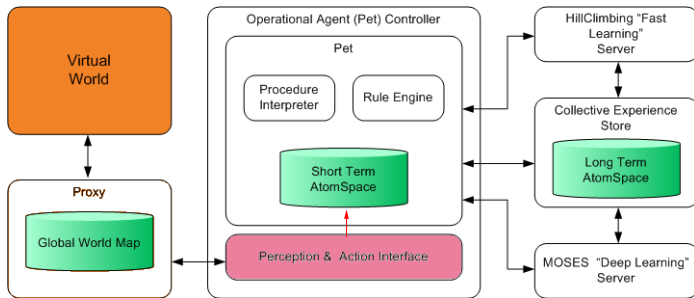


Proxy communicates with the OPC by the intermediate of the **Perception Action Interface**.

Outline

- 1 Introduction
- 2 Embodiment Sub-components**
 - Proxy
 - Operational Agent (Pet) Controller**
 - Learning Server
 - Collective Experience Store
- 3 Demo...
- 4 Conclusion

Perception Action Interface Function



- Convert Perceptions coming from the Proxy into **atoms** in the **agent's atomSpace**

Perception Action Interface, perception → atoms




⇒ Proxy ⇒ PAI ⇒

```
AtTimeLink
  TimeNode: "15:32:24.182"
  EvaluationLink
    PredicateNode: "actionDone"
    ListLink
      AvatarNode: "owner"
      Node: "grab"
      ListLink
        ObjectNode: "ball_id"
```


- 1 Pet watches the avatar grabbing a ball
- 2 Proxy sends perception to PAI
- 3 PAI writes the perception into the Pet's atomSpace

Embodiment AtomSpace Extensions, **Space and Time servers**



- Space Server: SpaceMap wrapped in atom, associative container for fast retrieval.
 - SpaceMap:
 - 2D Grid
 - Objects inside
 - Spatial functions and predicates, distance, near, ...

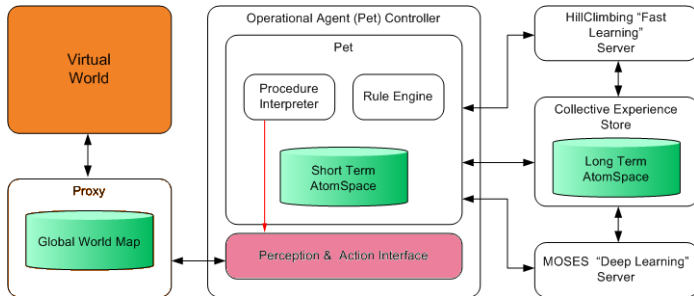
Embodiment AtomSpace Extensions, **Space and Time servers**



- Space Server: SpaceMap wrapped in atom, associative container for fast retrieval.
 - SpaceMap:
 - 2D Grid
 - Objects inside
 - Spatial functions and predicates, *distance*, *near*, ...
- Time Server: indexing timestamped atoms for fast retrieval

Episodic memory

Combo interpreter actions to action plan



- Convert actions coming from the **Combo interpreter** into **action plans** (sequence of actions) to the Proxy


Combo interpreter actions to action plan

```
goto_obj(barrel2)  
    PAI⇓  
{walk(546, 453), walk(547, 451), ...}  
    Proxy⇓
```



- 1 Combo interpreter + PAI create action plan
- 2 PAI sends it to Proxy, which sends it to virtual world


Rule Engine



Rule Engine Function

- Update agent's internal state. Example pet:
 - agent's feeling (hunger, happiness, ...)
 - agent's relation (familiar_with, friend_of, ...)

Rule Engine



Rule Engine Function

- Update agent's internal state. Example pet:
 - agent's feeling (hunger, happiness, ...)
 - agent's relation (familiar_with, friend_of, ...)
- Select next action to **satisfy its goals**. Example pet:
 - Fulfill needs (hunger, thirst, curiosity, etc)

Rule Engine, how it works?

For the moment it is really crude!

Rule Engine: Main loop

Evaluates a set of hard-coded rules to update state and choose next actions.

Rule Engine, how it works?

For the moment it is really crude!

Rule Engine: Main loop

Evaluates a set of hard-coded rules to update state and choose next actions.

3 types of rules:

1 Feeling rule

```
has_eaten(self)  $\Rightarrow$  happiness<0.8>
```

Rule Engine, how it works?

For the moment it is really crude!

Rule Engine: Main loop

Evaluates a set of hard-coded rules to update state and choose next actions.

3 types of rules:

① **Feeling rule**

`has_eaten(self) ⇒ happiness<0.8>`

② **Relation rule**

`has_licked(self X) ⇒ familiar_with(self X)<0.9>`

Rule Engine, how it works?

For the moment it is really crude!

Rule Engine: Main loop

Evaluates a set of hard-coded rules to update state and choose next actions.

3 types of rules:

1 **Feeling rule**

`has_eaten(self) \Rightarrow happiness<0.8>`

2 **Relation rule**

`has_licked(self X) \Rightarrow familiar_with(self X)<0.9>`

3 **Schema rule**

`is_hungry(self) $\xRightarrow{0.6}$ goto_find_food()`

Rule Engine, how it works?

For the petBrain we have **over 100 hard-coded rules!**


- 21 feeling rules
- 10 relation rules
- 76 schema rules

+ rule to use tricks learned by imitation

Outline

- 1 Introduction
- 2 Embodiment Sub-components**
 - Proxy
 - Operational Agent (Pet) Controller
 - Learning Server**
 - Collective Experience Store
- 3 Demo...
- 4 Conclusion

Learning Server




Lerning Server (LS) Function

Complex learning tasks (like imitation learning) run seperatly to not pertubate the agent's reactivity.

- Imitation Learning (Implemented)

Learning Server



Lerning Server (LS) Function

Complex learning tasks (like imitation learning) run seperatly to not pertubate the agent's reactivity.

- Imitation Learning (Implemented)
- Spontaneous Learning (not Implemented)

Imitation Learning



- 1 Pet's owner requests **learning session** and shows the trick

Imitation Learning



episodic memory
of learning session

Learning Server (LS)

- 1 Pet's owner requests **learning session** and shows the trick
- 2 OPC sends portion of atomSpace, **episodic memory of the learning session** to LS

Imitation Learning



episodic memory
of learning session

Learning Server (LS)
Search Combo that mimics
avatar's behavior

- 1 Pet's owner requests **learning session** and shows the trick
- 2 OPC sends portion of atomSpace, **episodic memory of the learning session to LS**
- 3 LS searches a **Combo program** that mimics that trick (using Hillclimbing or MOSES)

Imitation Learning




episodic memory
of learning session

Learning Server (LS)
Search Combo that mimics
avatar's behavior


- 1 Pet's owner requests **learning session** and shows the trick
- 2 OPC sends portion of atomSpace, **episodic memory of the learning session to LS**
- 3 LS searches a **Combo program** that mimics that trick (using Hillclimbing or MOSES)
- 4 Each Combo candidate is **run inside the imaginary world** of that episodic memory, the result is **compared with the avatar's behavior**

Imitation Learning



- 1 Owner asks the pet to **try the trick**

Imitation Learning



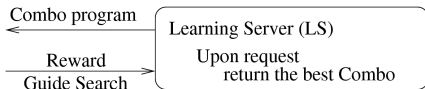

Combo program

Learning Server (LS)

Upon request
return the best Combo

- 1 Owner asks the pet to **try the trick**
- 2 LS sends its **best candidate so far**, and pet execute the trick

Imitation Learning



- 1 Owner asks the pet to **try the trick**
- 2 LS sends its **best candidate so far**, and pet execute the trick
- 3 Owner sends **positive or negative reward to guide** the search

Example of trick learning: Double dance

Dance in loop on cue

- Owner kicks left leg \Rightarrow tap dance
- Owner kicks right leg \Rightarrow lean rock dance


Combo to learn

```
boolean_while(not(says(owner "stop dancing"))  
              action_if(last_action(owner "kickL")  
                        tap_dance  
                        lean_rock_dance))
```

Outline

- 1 Introduction
- 2 Embodiment Sub-components**
 - Proxy
 - Operational Agent (Pet) Controller
 - Learning Server
 - Collective Experience Store**
- 3 Demo...
- 4 Conclusion

Collective Experience Store (not implemented yet)



Collective Consciousness

All agents can put in common their knowledge. When one gets smarter everybody takes advantage of it.

Outline

- 1 Introduction
- 2 Embodiment Sub-components
 - Proxy
 - Operational Agent (Pet) Controller
 - Learning Server
 - Collective Experience Store
- 3 Demo...**
- 4 Conclusion

Outline

- 1 Introduction
- 2 Embodiment Sub-components
 - Proxy
 - Operational Agent (Pet) Controller
 - Learning Server
 - Collective Experience Store
- 3 Demo...
- 4 Conclusion**

Conclusion

What has been done:

- Hard-coded behaviors
- Imitation learning
- but no spontaneous learning
- No language understanding yet (only commands)

Conclusion

What **remains to be done**:

- Spontaneous learning, concept creation, integrate Attention Allocation
- Transfer learning, Collective Experience Store
- Natural language processing and generation (in development)