

Automated Program Learning

MOSES

Nil Geisweiller

Novamente LLC

Xiamen University
AGI Summer School 2009

- 1 Introduction
- 2 Representation-Building
- 3 Optimization
- 4 Deme management
- 5 Demo...
- 6 Conclusion

Outline

- 1 Introduction
- 2 Representation-Building
- 3 Optimization
- 4 Deme management
- 5 Demo...
- 6 Conclusion

What is MOSES?

MOSES (*Meta-Optimizing Semantic Evolutionary Search*)



Evolutionary program learning, PhD Moshe Looks.



- 1 Search programs that **maximize the fitness** function

What is MOSES?

MOSES (*Meta-Optimizing Semantic Evolutionary Search*)



Evolutionary program learning, PhD Moshe Looks.



- 1 Search programs that **maximize the fitness** function
- 2 Take advantage of **program semantics** and program space topology

What is MOSES?

MOSES (*Meta-Optimizing Semantic Evolutionary Search*)

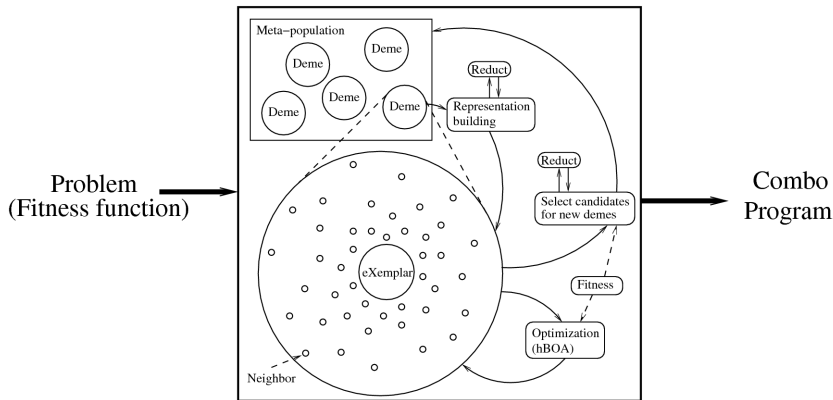


Evolutionary program learning, PhD Moshe Looks.



- 1 Search programs that **maximize the fitness** function
- 2 Take advantage of **program semantics** and program space topology
- 3 Attempt to discover **fitness landscape regularities to speed up the search**

How it works?



How it works?

- 1 Reduction in **normal form**
 - minimize over-representation
 - improve syntactic vs semantic distance correlation
 - simplify or even improve structure

How it works?

- 1 Reduction in **normal form**
 - minimize over-representation
 - improve syntactic vs semantic distance correlation
 - simplify or even improve structure
- 2 Population building, **representation-building** defines the **deme's neighborhood**

How it works?

- 1 Reduction in **normal form**
 - minimize over-representation
 - improve syntactic vs semantic distance correlation
 - simplify or even improve structure
- 2 Population building, **representation-building** defines the **deme's neighborhood**
- 3 Optimization, find the best candidates in the deme's neighborhood. Learn how to **differentiate good vs bad programs** and **bias the search** accordingly
 - hBOA
 - Building-Block Hill Climbing (under development)

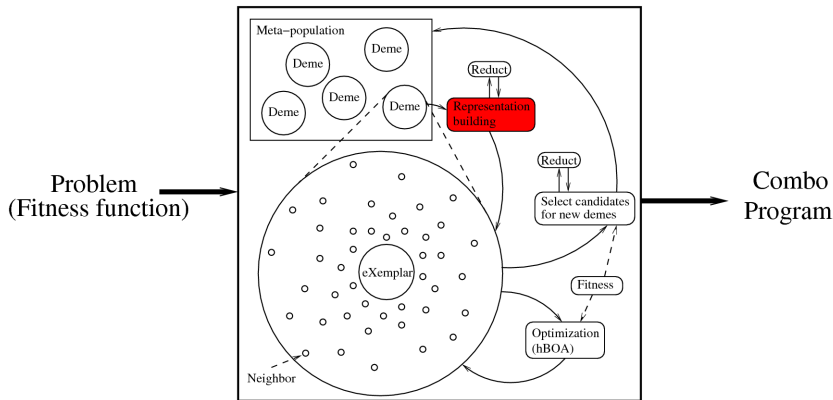
How it works?

- 1 Reduction in **normal form**
 - minimize over-representation
 - improve syntactic vs semantic distance correlation
 - simplify or even improve structure
- 2 Population building, **representation-building** defines the **deme's neighborhood**
- 3 Optimization, find the best candidates in the deme's neighborhood. Learn how to **differentiate good vs bad programs** and **bias the search** accordingly
 - hBOA
 - Building-Block Hill Climbing (under development)
- 4 Deme management
 - Set of demes, **meta-population**
 - Diversity, **preserving interesting demes**

Outline

- 1 Introduction
- 2 Representation-Building**
- 3 Optimization
- 4 Deme management
- 5 Demo...
- 6 Conclusion

Representation-building: Build a deme's population



Building string of knobs

Population of a Deme

Centered around an exemplar, each neighbor is a variation of that exemplar according to the representation-building, a **string of knobs**.

`and(#1 not(#2))`

\Rightarrow



Building string of knobs

Domain specific rules to create knobs, example in the Boolean domain with `and(x not(y))`:



Building string of knobs

Domain specific rules to create knobs, example in the Boolean domain with `and(x not(y))`:

- 1 Under every junctor, $\forall v$ **not already sibling**, add $[\emptyset, v, \text{not}(v)]$



Building string of knobs

Domain specific rules to create knobs, example in the Boolean domain with `and(x not(y))`:


- 1 Under every junctor, $\forall v$ **not already sibling**, add $[\emptyset, v, \text{not}(v)]$
- 2 Any junctor can be flipped



Building string of knobs

Domain specific rules to create knobs, example in the Boolean domain with $\text{and}(x \text{ not } (y))$:


- 1 Under every junctor, $\forall v$ **not already sibling**, add $[\emptyset, v, \text{not}(v)]$
- 2 Any junctor can be flipped
- 3 Under every junctor add **opposite junctor** + children



Building string of knobs

Domain specific rules to create knobs, example in the Boolean domain with $\text{and}(x \text{ not } (y))$:


- 1 Under every junctor, $\forall v$ **not already sibling**, add $[\emptyset, v, \text{not}(v)]$
- 2 Any junctor can be flipped
- 3 Under every junctor add **opposite junctor** + children
- 4 Insert an **opposite junctor** above the root + children



Building string of knobs

Domain specific rules to create knobs, example in the Boolean domain with $\text{and}(x \text{ not } (y))$:

- 1 Under every junctor, $\forall v$ **not already sibling**, add $[\emptyset, v, \text{not}(v)]$
- 2 Any junctor can be flipped
- 3 Under every junctor add **opposite junctor** + children
- 4 Insert an **opposite junctor** above the root + children
- 5 And a few more...



Building string of knobs




Figure: Built knobs for $and(x_1 not(x_4))$, extracted from Moshe's PhD.

Optimization problem takes place on the knob string



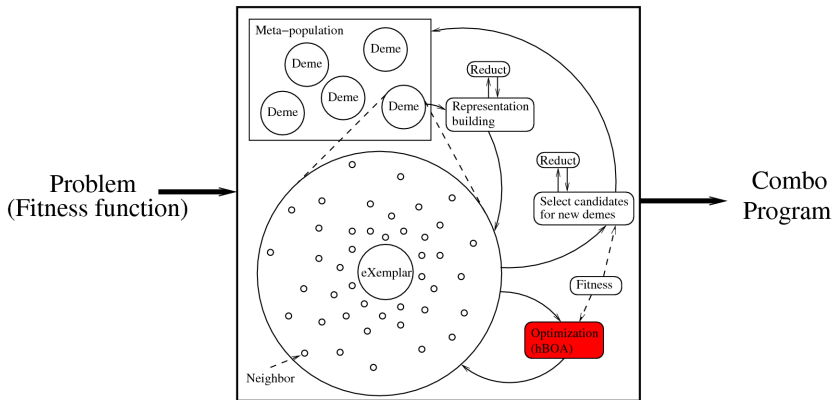
For example:

knob setting	combo (reduced)	distance
(and, \emptyset , \emptyset , \emptyset , \emptyset)	and(x not(y))	0
(or, \emptyset , \emptyset , \emptyset , \emptyset)	or(x not(y))	1
(and, \emptyset , \emptyset , not(x), \emptyset)	or(not(x) not(y))	1
(or, \emptyset , \emptyset , not(x), \emptyset)	true	2

Outline

- 1 Introduction
- 2 Representation-Building
- 3 Optimization**
- 4 Deme management
- 5 Demo...
- 6 Conclusion

Optimization: Find the best candidates inside a deme



MOSES' Optimization algorithms

- 1 hBOA, multivariate model-building (not yet ported to the OpenCog version, univariate model-building instead)

MOSES' Optimization algorithms


- 1 hBOA, multivariate model-building (not yet ported to the OpenCog version, univariate model-building instead)
- 2 Hill-Climbing

MOSES' Optimization algorithms

- 1 hBOA, multivariate model-building (not yet ported to the OpenCog version, univariate model-building instead)
- 2 Hill-Climbing
- 3 Building-Block Hill-Climbing (being ported to the OpenCog version)


Hierarchical Bayesian Optimization Algorithm (hBOA)

<i>Candidate</i>	<i>score</i>
00001010	0.2
00011010	0.6
01000100	0.01
00110110	0.5
10010010	0.6
⋮	⋮



Hierarchical Bayesian Optimization Algorithm (hBOA)


<i>Candidate</i>	<i>score</i>
00001010	0.2
00011010	0.6
01000100	0.01
00110110	0.5
10010010	0.6
⋮	⋮



- Split the population in **good vs bad candidates**

Hierarchical Bayesian Optimization Algorithm (hBOA)

<i>Candidate</i>	<i>score</i>
00001010	0.2
00011010	0.6
01000100	0.01
00110110	0.5
10010010	0.6
⋮	⋮



- Split the population in **good vs bad candidates**
- **Learn a classifier**, but not too strict or incorrect

Hierarchical Bayesian Optimization Algorithm (hBOA)

Candidate	score
00001010	0.2
00011010	0.6
01000100	0.01
00110110	0.5
10010010	0.6
⋮	⋮



- Split the population in **good vs bad candidates**
- **Learn a classifier**, but not too strict or incorrect
- Probabilistic classifier, **distribution of good candidates**

Hierarchical Bayesian Optimization Algorithm (hBOA)


<i>Candidate</i>	<i>score</i>
00001010	0.2
00011010	0.6
01000100	0.01
00110110	0.5
10010010	0.6
⋮	⋮



- Split the population in **good vs bad candidates**
- **Learn a classifier**, but not too strict or incorrect
- Probabilistic classifier, **distribution of good candidates**
- **Sample new candidates** according to the distribution

Hierarchical Bayesian Optimization Algorithm (hBOA)


<i>Candidate</i>	<i>score</i>
10011010	0.5
00111011	0.6
11111001	0.7
00111010	0.4
10000001	0.1
⋮	⋮





- Split the population in **good vs bad candidates**
- **Learn a classifier**, but not too strict or incorrect
- Probabilistic classifier, **distribution of good candidates**
- **Sample new candidates** according to the distribution
- Repeat on the new population

Distribution of Good Candidates: Bayesian Network

<i>Candidate (X_i)</i> <i>i : 12345678</i>	<i>score</i>
00001010	0.2
00011010	0.6
01000100	0.01
00110110	0.5
10010010	0.6
⋮	⋮

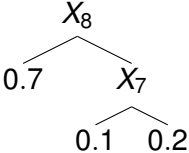


Decomposing the problem into sub-problems



Conditional Probability with Decision Tree

<i>Marginal</i>	<i>Prob</i>
$P(X_1 = 1)$	0.3
$P(X_2 = 1)$	0.05
$P(X_4 = 1)$	0.9
$P(X_7 = 1)$	0.88

<i>Conditional</i>	<i>Prob</i>
$P(X_8 X_7)$	
$P(X_5 X_1)$...
$P(X_6 X_1, X_5)$...
$P(X_3 X_6)$...

In OpenCog for the moment only **univariate**

Univariate


Only marginal probabilities

<i>Marginal</i>	<i>Prob</i>
$P(X_1 = 1)$	0.3
$P(X_2 = 1)$	0.05
$P(X_3 = 1)$	0.2
$P(X_4 = 1)$	0.5
$P(X_5 = 1)$	0.4
$P(X_6 = 1)$	0.7
$P(X_7 = 1)$	0.88
$P(X_8 = 1)$	0.1

Hill-Climbing, Building-Block Hill-Climbing

- Hill-Climbing


The problem with hill climbing is that it gets stuck on "local-maxima"



Hill-Climbing, Building-Block Hill-Climbing

- Hill-Climbing

The problem with hill climbing is that it gets stuck on "local-maxima"

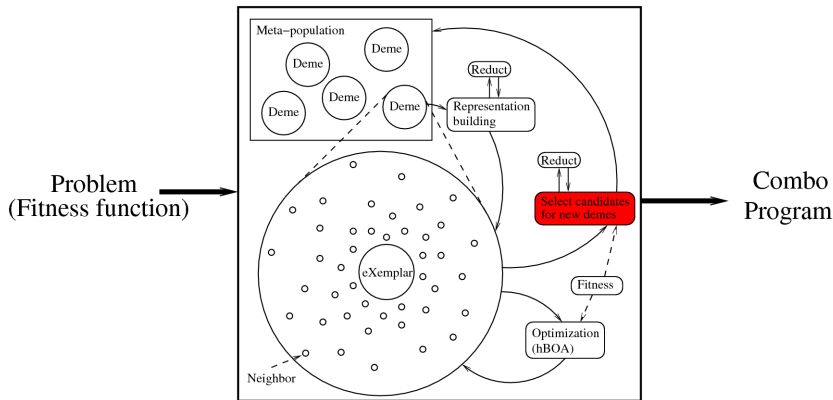


- Building-Block Hill-Climbing \Rightarrow Redefine neighborhood to take short-cuts.

Outline

- 1 Introduction
- 2 Representation-Building
- 3 Optimization
- 4 Deme management**
- 5 Demo...
- 6 Conclusion

Select Candidates for Future Demes



Preserving Diversity, candidates that behave differently and non-dominated

Preserving Diversity, candidates that behave differently and non-dominated



Preserving Diversity, candidates that behave differently and non-dominated



Preserving Diversity, candidates that behave differently and non-dominated



- Neither one dominates the other

Preserving Diversity, candidates that behave differently and non-dominated



- Neither one dominates the other
- But “Moshe Lewis” dominates then both

Behavioral score

Behavioral score

Partial order

- if $c1 < c2$ then **$c1$ is dominated by $c2$**

Behavioral score

Behavioral score

Partial order

- if $c1 < c2$ then **$c1$ is dominated by $c2$**
- if $c1 > c2$ then **$c1$ dominates $c2$**

Behavioral score

Behavioral score

Partial order

- if $c1 < c2$ then **$c1$ is dominated by $c2$**
- if $c1 > c2$ then **$c1$ dominates $c2$**
- otherwise, **neither one dominates the other**

Behavioral score

Behavioral score

Partial order

- if $c1 < c2$ then **$c1$ is dominated by $c2$**
- if $c1 > c2$ then **$c1$ dominates $c2$**
- otherwise, **neither one dominates the other**

Behavioral score

Behavioral score

Partial order

- if $c1 < c2$ then **$c1$ is dominated by $c2$**
- if $c1 > c2$ then **$c1$ dominates $c2$**
- otherwise, **neither one dominates the other**

For instance: **vector of floats** (f_1, \dots, f_n) where each f_i measure how well a candidate is doing for that particular feature.

If for all features i , $c1$ is doing better than $c2$, then $c1$ dominates $c2$.


Deme selection

Keep non-dominated exemplars as potential deme.

Deme selection

Keep non-dominated exemplars as potential deme.


We want to **learn** $f(x)$.



Deme selection

Keep non-dominated exemplars as potential deme.


We want to learn $f(x)$.



Deme selection

Keep non-dominated exemplars as potential deme.


We want to **learn** $f(x)$.



Deme selection

Keep non-dominated exemplars as potential deme.


We want to learn $f(x)$.



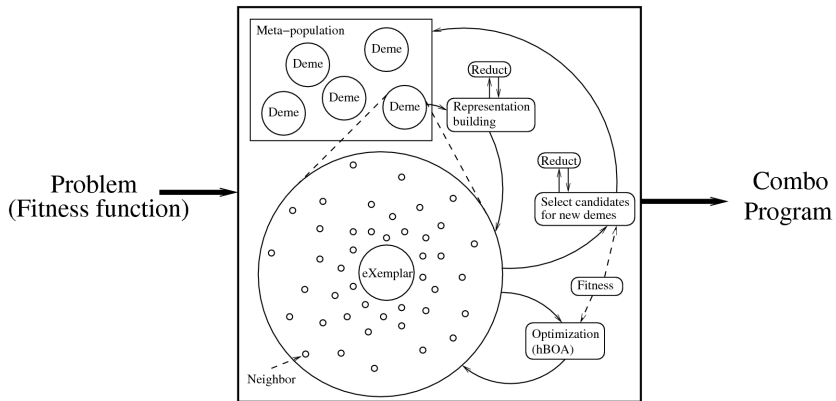
Deme selection

Keep non-dominated exemplars as potential deme.

We want to learn $f(x)$.



$k(x)$ dominates $g(x)$ and will not be included in the meta-population



Some benchmark

Technique	Computational Effort			
	3-parity	4-parity	5-parity	6-parity
Univariate MOSES	6,151	73,977	2,402,523	342,280,092
Evolutionary programming [15]	28,500	181,500	2,100,000	no solutions
Genetic programming [49]	96,000	384,000	6,528,000	no solutions
MOSES	5,112	72,384	1,581,212	100,490,013

Figure: Computational effort to find solution of n-parity 99% of the time (*extracted from Moshe's PhD thesis*)

$$n\text{-parity}(b_1, \dots, b_n) = \text{even} \left(\sum_{i=1}^n \text{int}(b_i) \right)$$

Outline

- 1 Introduction
- 2 Representation-Building
- 3 Optimization
- 4 Deme management
- 5 Demo...**
- 6 Conclusion

Outline

- 1 Introduction
- 2 Representation-Building
- 3 Optimization
- 4 Deme management
- 5 Demo...
- 6 Conclusion**

Conclusion

MOSES outperforms Genetic Algorithms because:

- 1 minimize over-representation (**Reduction in normal form**)

Conclusion

MOSES outperforms Genetic Algorithms because:

- 1 minimize over-representation (**Reduction in normal form**)
- 2 build expressive deme population by taking into account operator semantics (**representation-building**)

Conclusion

MOSES outperforms Genetic Algorithms because:

- 1 minimize over-representation (**Reduction in normal form**)
- 2 build expressive deme population by taking into account operator semantics (**representation-building**)
- 3 Maintain diversity in the meta-population (**deme managment**)

Conclusion

MOSES outperforms Genetic Algorithms because:

- 1 minimize over-representation (**Reduction in normal form**)
- 2 build expressive deme population by taking into account operator semantics (**representation-building**)
- 3 Maintain diversity in the meta-population (**deme managment**)
- 4 Attempt to find regularities in one deme's population to speed up optimization (**model-building**)

What remains to be done?

- 1 Only handles Boolean, continuous, numerico-boolean and (partially) action-perception expressions
⇒ more programmatic construct

What remains to be done?

- 1 Only handles Boolean, continuous, numerico-boolean and (partially) action-perception expressions
⇒ **more programmatic construct**
- 2 Representation-building is hard-coded per operator set
⇒ Generalized for **operator properties (especially useful if combined with PLEASURE)**

What remains to be done?

- 1 Only handles Boolean, continuous, numerico-boolean and (partially) action-perception expressions
⇒ **more programmatic construct**
- 2 Representation-building is hard-coded per operator set
⇒ Generalized for **operator properties (especially useful if combined with PLEASURE)**
- 3 Model-building is slow
⇒ **improve model-building** (better Bayesian learning, transfer learning across populations), or use other optimization method

What remains to be done?

- 1 Only handles Boolean, continuous, numerico-boolean and (partially) action-perception expressions
⇒ **more programmatic construct**
- 2 Representation-building is hard-coded per operator set
⇒ Generalized for **operator properties (especially useful if combined with PLEASURE)**
- 3 Model-building is slow
⇒ **improve model-building** (better Bayesian learning, transfer learning across populations), or use other optimization method
- 4 No transfer learning across problem instances
⇒ Integrative AGI, Attention Allocation, PLN, etc.